



香港中文大學 (深圳)

The Chinese University of Hong Kong

# CSC3100 Data Structures

## Lecture 14: Binary search tree

Li Jiang

School of Data Science (SDS)

The Chinese University of Hong Kong, Shenzhen

---



# Outline




---

- ▶ In this lecture, we will learn
  - Binary search tree (BST)
  - Operations on BST
    - Search a key
    - Find the minimum/maximum and find successor/predecessor
    - Insert and delete
  - Exercises



# Binary search tree (BST) property

- ▶ BST is a binary tree such that for each node  $T$ ,
  - the key values in its left subtree are **smaller** than the key value of  $T$
  - the key values in its right subtree are **larger** than the key value of  $T$





# Applications of BST

---

- ▶ Many applications due to its **ordered structure**
  - Useful for indexing and multi-level indexing
  - Helpful in maintaining a sorted stream of data
  - Helpful to implement various searching algorithms and data structures (e.g., TreeMap, TreeSet, Priority queue)

java.util

**Class TreeMap<K,V>**

java.lang.Object

java.util.AbstractMap<K,V>

java.util.TreeMap<K,V>

java.util

**Class TreeSet<E>**

java.lang.Object

java.util.AbstractCollection<E>

java.util.AbstractSet<E>

java.util.TreeSet<E>



# BST ADT

---


- ▶ Support many dynamic set operations
  - searchKey, findMin, findMax, predecessor, successor, insert, delete
- ▶ Running time of basic operations on BST
  - On average:  $\Theta(\log n)$ 
    - The expected height of the tree is  $\log n$
  - In the worst case:  $\Theta(n)$ 
    - The tree is a linear chain of  $n$  nodes



# Searching for a key

- ▶ Given a pointer to the root of a tree and a key  $k$ :
  - Return a pointer to a node with key  $k$  if one exists, otherwise return NIL

- ▶ Example




- ▶ Search for key 13:
  - $15 \rightarrow 6 \rightarrow 7 \rightarrow 13$



# Searching for a key

find(x, k)

1. **if**  $x = \text{NIL}$  or  $k = \text{key}[x]$
2.     **return**  $x$
3. **if**  $k < \text{key}[x]$
4.     **return** find(left [x], k )
5. **else**
6.     **return** find(right [x], k )



Running time:  $O(h)$ , where  $h$  is the height of tree




# Finding the minimum

- ▶ Goal: find the minimum value in a BST
  - Following **left child pointers** from the root, until a NIL is encountered

findMin(x)

1. **while** left [x]  $\neq$  NIL
2.       **do** x  $\leftarrow$  left [x]
3. **return** x



Minimum = 2

Running time:  $O(h)$ , where  $h$  is the height of tree






# Finding the maximum

- ▶ Goal: find the maximum value in a BST
  - Following **right child pointers** from the root, until a NIL is encountered

findMax(x)

1. **while** right [x]  $\neq$  NIL
2.       **do**  $x \leftarrow$  right [x]
3. **return** x



Maximum = 20


Running time:  $O(h)$ , where  $h$  is the height of tree



# Successor

**Def:**  $\text{successor}(x) = y$ , such that  $\text{key}[y]$  is the smallest key  $> \text{key}[x]$

- ▶ **E.g.:**  $\text{successor}(15) = 17$   
 $\text{successor}(13) = 15$   
 $\text{successor}(9) = 13$




- ▶ **Case 1: right (x) is non-empty**
  - $\text{successor}(x)$  = the minimum in right (x)
- ▶ **Case 2: right (x) is empty**
  - go up the tree until the current node is a left child:  $\text{successor}(x)$  is the parent of the current node
  - if you cannot go further (and you reached the root): x is the largest element



# Successor

successor(x)

1. **if** right [x]  $\neq$  NIL
2.     **return** findMin(right [x])
3.  $y \leftarrow p[x]$
4. **while**  $y \neq \text{NIL}$  and  $x = \text{right } [y]$
5.     **do**  $x \leftarrow y$
6.      $y \leftarrow p[y]$
7. **return** y




Running time:  $O(h)$ , where  $h$  is the height of tree



# Predecessor

**Def:** predecessor ( $x$ ) =  $y$ , such that key [ $y$ ] is the biggest key  $<$  key [ $x$ ]

- ▶ **E.g.:** predecessor (15) = 13  
predecessor (9) = 7  
predecessor (7) = 6



- ▶ Case 1: left ( $x$ ) is non-empty
  - predecessor ( $x$ ) = the maximum in left ( $x$ )
- ▶ Case 2: left ( $x$ ) is empty
  - go up the tree until the current node is a right child: predecessor ( $x$ ) is the parent of the current node
  - if you cannot go further (and you reached the root):  $x$  is the smallest element



# Insertion





---

- ▶ Goal: Insert value  $v$  into a binary search tree
  
- ▶ Find the position and insert as a leaf:
  - If  $\text{key}[x] < v$  move to the right child of  $x$ ,  
else move to the left child of  $x$
  - When  $x$  is NIL, we found the correct position
  - If  $v < \text{key}[y]$  insert the new node as  $y$ 's left child  
else insert it as  $y$ 's right child
  
- Beginning at the root, go down the tree and maintain:
  - Pointer  $x$ : traces the downward path (current node)
  - Pointer  $y$ : parent of  $x$  ("trailing pointer" )



# Example


Insert 13:





# Insert algorithm

```
1.  y ← NIL
2.  x ← root [T]
3.  while x ≠ NIL
4.    do y ← x
5.    if key [z] < key [x]      z: the node to
6.      x ← left [x]           be inserted
7.    else
8.      x ← right [x]
9.  p[z] ← y
10. if y = NIL
11.   root [T] ← z    ▷ T was empty
12. else
13.   if key [z] < key [y]
14.     left [y] ← z
15.   else
16.     right [y] ← z
```



Best-case and worst-case time complexities?

Running time:  $O(h)$



# Exercise

---



- ▶ Build a binary search tree for the following sequence  
15, 6, 18, 3, 7, 17, 20, 2, 4





# Deletion

- ▶ Goal: Delete a given node  $z$  from a binary search tree
- ▶ Idea:
  - **Case 1:**  $z$  has no children
    - Delete  $z$  by making the parent of  $z$  point to NIL







# Deletion

## ► Case 2: z has one child

- Delete z by making the parent of z point to z's child, instead of to z, and link the parent with the new child






# Deletion

We can also replace it by "Find z's predecessor y (rightmost node in z's left subtree)"

## ► Case 3: z has two children


- Find z's **successor** y (leftmost node in z's right subtree)
- y has either no or one right child (but no left child), why?
- Delete y from the tree (via Case 1 or 2)
- Replace z's key by y's key, and satellite data with y's





# Deletion algorithm


1. **if**  $\text{left}[z] = \text{NIL}$  and  $\text{right}[z] = \text{NIL}$  //z has no child
2.     **if**  $p[z] = \text{NIL}$  **then**  $\text{root}[T] = \text{NIL}$
3.     **else**
4.         **if**  $z = \text{left}[p[z]]$
5.              $\text{left}[p[z]] = \text{NIL}$
6.         **else**
7.              $\text{right}[p[z]] = \text{NIL}$





# Deletion algorithm


```
1.  if left[z] = NIL and right[z] ≠ NIL  //z has one right child
2.      y = right[z]
3.      if p[z] = NIL
4.          root[T] = y
5.      else
6.          p[y] = p[z]
7.          if z = left[p[z]]
8.              left[p[z]] = y
9.          else
10.             right[p[z]] = y
11. if left[z] ≠ NIL and right[z] = NIL  //z has one left child
12.     y = left[z]
13.     if p[z] = NIL
14.         root[T] = y
15.     else
16.         p[y] = p[z]
17.         if z = left[p[z]]
18.             left[p[z]] = y
19.         else
20.             right[p[z]] = y
```





# Deletion algorithm

```
1.  if left[z] ≠ NIL and right[z] ≠ NIL //z has two children
2.    y ← TREE-SUCCESSOR(z)      //left-most node in right tree
3.    if p[y] = z
4.      right[z] = right[y]
5.      if right[y] ≠ NIL
6.        p[right[y]] = z
7.    else
8.      if right[y] = NIL
9.        left[p[y]] ← NIL
10.     else
11.       x ← right[y]
12.       p[x] ← p[y]
13.       left[p[y]] ← x
14.  key[z] ← key[y] //copy y's data into z
```



Best/worst-case time complexities?



# Summary

---

- ▶ Operations on binary search trees:
  - Search  $O(h)$
  - Predecessor  $O(h)$
  - Successor  $O(h)$
  - FindMin  $O(h)$
  - FindMax  $O(h)$
  - Insert/Delete  $O(h)$
  
- ▶ These operations are fast if the height of the tree is **small** - otherwise their performance is similar to that of a linked list



# Binary search trees vs linear lists

---

Operation	BST	Sorted- array-based List	Linked List
Constructor	$O(1)$	$O(1)$	$O(1)$
IsFull	$O(1)$	$O(1)$	$O(1)$
IsEmpty	$O(1)$	$O(1)$	$O(1)$
RetrieveItem	$O(\log N)^*$	$O(\log N)$	$O(N)$
InsertItem	$O(\log N)^*$	$O(N)$	$O(N)$
DeleteItem	$O(\log N)^*$	$O(N)$	$O(N)$

\*assuming  $h = O(\log N)$





# The issues in BST


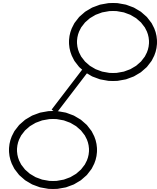
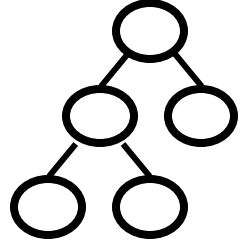
---

- ▶ After a series of delete operations, the above algorithm favors making **the left sub-trees deeper than the right**
- ▶ One solution:
  - Try to eliminate the problem by randomly choosing between the smallest element in the right sub-tree and the largest in the left when replacing the deleted element (not rigorous and not prove it yet!!)
- ▶ Existing balanced BST solutions
  - AVL tree: height  $O(\log n)$
  - Red-black tree: height  $O(\log n)$



# Exercise 1: count leaves

Example:

<p>A NULL binary tree has <b>0</b> leaf node</p>	 <p>A tree with 1 node has <b>1</b> leaf node</p>	 <p>No. of leaf nodes = <b>1</b></p>	 <p>No. of leaf nodes = <b>3</b></p>
--------------------------------------------------	------------------------------------------------------------------------------------------------------------------------------------	-------------------------------------------------------------------------------------------------------------------------	-------------------------------------------------------------------------------------------------------------------------




//To count the number of leaf nodes

```
int Mytree::count_leaf(TreeNode* p)
{
    if (p == NULL)
        return 0;
    else if ((p->left == NULL) && (p->right == NULL))
        return 1;
    else
        return count_leaf(p->left) + count_leaf(p->right);
}
```



## Exercise 2: operation commutative

- ▶ In a binary search tree, are the insert and delete operations commutative?
  - $\text{delete}(a)$  then  $\text{delete}(b) \Leftrightarrow \text{delete}(b)$  then  $\text{delete}(a)$ ?
  - $\text{insert}(a)$  then  $\text{insert}(b) \Leftrightarrow \text{insert}(b)$  then  $\text{insert}(a)$ ?



Case 1: Delete 5 and then 6

Case 2: Delete 6 and then 5



## Exercise 3: sorting with BST

---

- ▶ How to sort an array of keys by building and traversing a BST?

```
1. Sort (A[ ])
2.   initialize a BST T
3.   for i = 1 to n
4.       insert(A[i]) into T
5.   inorder-tree-walk(T)
```


- What are the worst case and best case time costs?
- In practice, how would this compare to other sorting algorithms?



# Exercise 4: lowest common ancestor

## ► Lowest common ancestor (LCA):

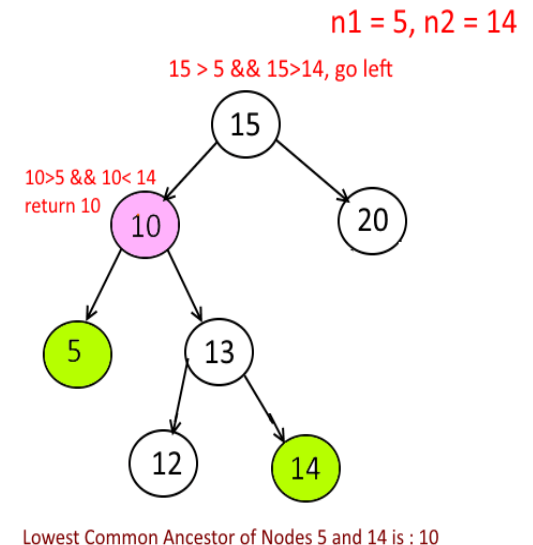
- The LCA of two nodes  $n1$  and  $n2$  is a node  $X$  such that node  $X$  will be the lowest node who has both  $n1$  and  $n2$  as its descendants
- Given a BST and two nodes  $n1$  and  $n2$ , how to find their LCA?



Lowest Ancestor Ancestor (5, 21) = 20  
Lowest Ancestor Ancestor (10, 30) = 25  
Lowest Ancestor Ancestor (5, 32) = 25  
Lowest Ancestor Ancestor (10, 23) = 20

Approach:

- 1) Start with the root
- 2) If  $root > n1$  and  $root > n2$  then lowest common ancestor will be in left subtree
- 3) If  $root < n1$  and  $root < n2$  then lowest common ancestor will be in right subtree
- 4) If Step 2 and Step 3 is false then we are at the root which is LCA, return it





# Recommended reading

---

- ▶ Reading this week
  - Chapter 12, textbook
- ▶ Next lecture
  - AVL-tree: Chapter 12, textbook