



香港中文大學 (深圳)
The Chinese University of Hong Kong

(Materials of this lecture are NOT included in the midterm and final exams)

CSC3100 Data Structures

Lecture 16: Red-black tree

Li Jiang
School of Data Science (SDS)
The Chinese University of Hong Kong, Shenzhen



Outline

- ▶ Definitions and examples
- ▶ Properties and the height of the tree
- ▶ Operations
 - Insertion algorithm with three cases
 - Deletion algorithm (homework)



Red-black tree

- ▶ A “balanced” binary search tree
 - A binary search tree has an additional attribute for its nodes: color which can be either **red** or **black**
 - It restricts the way that nodes can be colored on any path from the root to a leaf
 - It ensures that no path is more than twice as long as any other path
 - It guarantees an $O(\log n)$ running time for many operations, such as search, insertion, and deletion




Red-black tree properties

1. Every node is either red or black
2. The root is black
3. Every leaf (NIL) is black
4. If a node is red, then both its children are black
No two consecutive red nodes on a simple path from the root to a leaf
5. For each node, all paths from that node to descendant leaves contain the same number of black nodes




Example



- ▶ For convenience we use a sentinel $NIL[T]$ to represent all the NIL nodes at the leaves
 - $NIL[T]$ has the same fields as an ordinary node
 - $Color[NIL[T]] = BLACK$
 - The other fields may be set to arbitrary values



Black height of a node



- ▶ Height of a node x :
 - $h(x)$ is the number of edges in the longest path to a leaf
- ▶ Black-height of a node x :
 - $bh(x)$ is the number of black nodes (including NIL) on the path from x to a leaf, not counting x



Important property of red-black tree


A red-black tree with n internal nodes
has height at most $2\log(n + 1)$

- ▶ Need to prove two claims first ...



Claim 1


- ▶ Any node x with height $h(x)$ has $bh(x) \geq h(x)/2$
- ▶ **Proof**
 - By property 4, at most $h/2$ **red** nodes on the path from the node to a leaf
 - Hence at least $h/2$ are **black**





Claim 2

- ▶ The subtree rooted at any node x contains at least $2^{bh(x)} - 1$ internal nodes





Claim 2 (Cont'd)

Proof: By induction on $h[x]$

Basis: $h[x] = 0 \Rightarrow$

x is a leaf ($NIL[T]$) \Rightarrow

$bh(x) = 0 \Rightarrow$

of internal nodes: $2^0 - 1 = 0$


Inductive hypothesis: assume it is true for $h[x] = h-1$



Claim 2 (Cont'd)

Inductive step:

- ▶ Prove it for $h[x] = h$
- ▶ Let $bh(x) = b$. Then, any child y of x has:
 - $bh(y) = b$ (if the child is **red**), or
 - $bh(y) = b - 1$ (if the child is **black**)





Claim 2 (Cont'd)

- ▶ Using inductive hypothesis, the number of internal nodes for each child of x is at least (if it is black):


$$2^{bh(x) - 1} - 1$$

- ▶ The subtree rooted at x has at least:

$$(2^{bh(x) - 1} - 1) + (2^{bh(x) - 1} - 1) + 1$$

$$= 2 \cdot (2^{bh(x) - 1} - 1) + 1$$

$$= 2^{bh(x)} - 1 \text{ internal nodes}$$



$$bh(l) \geq bh(x) - 1$$

$$bh(r) \geq bh(x) - 1$$



Important property of red-black tree

A red-black tree with n internal nodes
has height at most $2\log(n + 1)$
Proof in the next slides.

- ▶ Claim 1: Any node x with height $h(x)$ has $bh(x) \geq h(x)/2$
- ▶ Claim 2: The subtree rooted at any node x contains **at least** $2^{bh(x)} - 1$ internal nodes



Height of red-black tree


Lemma: A red-black tree with n internal nodes has height at most $2\log(n + 1)$.

Proof:

$$n \geq 2^{bh} - 1 \geq 2^{h/2} - 1$$

number n
of internal
nodes

since $bh \geq h/2$



► Add 1 to both sides and then take logs:

$$n + 1 \geq 2^{bh} \geq 2^{h/2}$$

$$\log(n + 1) \geq h/2$$

$$\Rightarrow h \leq 2 \log(n + 1)$$



Operations on red-black tree

- ▶ The non-modifying operations: **MINIMUM**, **MAXIMUM**, and **SEARCH** run in $O(h)$ time
 - They take $O(\log n)$ time on red-black trees
 - **SEARCH** is similar to the search on binary search tree
- ▶ What about **INSERT** and **DELETE**?
 - We have to guarantee that the modified tree will still be a red-black tree
 - Reconstruction will be too expensive
 - They can still be completed in $O(\log n)$ time



INSERT operation


INSERT: Suppose we want to insert 35. What color to make the new node?

▶ Red?

- Property 4 is violated: if a node is red, then its children are black

▶ Black?

- Property 5 is violated: all paths from a node to its leaves contain the same number of black nodes






Rotations



- ▶ After insertion and deletion on red-black trees, we need to restore the red-black tree properties
- ▶ Rotations take a red-black tree and a node within the tree and:
 - Two types of rotations: Left & right rotations
 - Together with some node re-coloring they help restore the red-black tree property
 - Change some of the pointer structure
 - Do not change the binary search tree property




Left and right rotations



LEFT-ROTATE(T, x)
.....>>>



RIGHT-ROTATE(T, y)
.....>>>





INSERT

- ▶ Goal:
 - Insert a new node z into a red-black tree

- ▶ Idea:
 - Insert node z into the tree as for an ordinary BST
 - Color the node **red**
 - Restore the red-black tree properties
 - Use an auxiliary procedure **RB-INSERT-FIXUP**




Properties affected by INSERT

1. Every node is either red or black OK!
2. The root is black If the root is changed
⇒ May not OK
3. Every leaf (NIL) is black OK!
4. If a node is red, then both its children are black

If $p(z)$ is red ⇒ not OK
z and $p(z)$ are both red

- OK!
5. For each node, all paths from the node to descendant leaves contain the same number of black nodes








INSERT(T, z)

1. $y \leftarrow \text{NIL}$
 2. $x \leftarrow \text{root}[T]$
 3. **while** $x \neq \text{NIL}$
 4. **do** $y \leftarrow x$
 5. **if** $\text{key}[z] < \text{key}[x]$
 6. **then** $x \leftarrow \text{left}[x]$
 7. **else** $x \leftarrow \text{right}[x]$
 8. $p[z] \leftarrow y$
 9. **if** $y = \text{NIL}$
 10. **then** $\text{root}[T] \leftarrow z$
 11. **else if** $\text{key}[z] < \text{key}[y]$
 12. **then** $\text{left}[y] \leftarrow z$
 13. **else** $\text{right}[y] \leftarrow z$
 14. $\text{left}[z] \leftarrow \text{NIL}$
 15. $\text{right}[z] \leftarrow \text{NIL}$
 16. $\text{color}[z] \leftarrow \text{RED}$
 17. $\text{RB-INSERT-FIXUP}(T, z)$
- Initialize nodes x and y
- Throughout the algorithm y points to the parent of x
- Go down the tree until reaching a leaf
- At that point y is the parent of the node to be inserted
- Sets the parent of z to be y
- The tree was empty: set the new node to be the root
- Otherwise, set z to be the left or right child of y , depending on whether the inserted node is smaller or larger than y 's key
- Set the fields of the newly added node
- Fix any inconsistencies that could have been introduced by adding this new red node



RB-Insert-Fixup(T, z)

- ▶ Case 1: z's uncle y is red
 - Solution: recolor
- ▶ Case 2: z's uncle y is black and z is a right child
 - Solution: double rotation
 - Can be transferred to Case 3
- ▶ Case 3: z's uncle y is black and z is a left child
 - Solution: single rotation





RB-Insert-Fixup(T, z)

1. **while** z.p.color == red ← The while loop repeats only when
Case 1 is executed: $O(\log n)$ times

2. **if** z.p == z.p.p.left

3. y = z.p.p.right

4. **if** y.color == red

5. z.p.color = black // case 1

6. y.color = black // case 1

7. z.p.p.color = red // case 1

8. z = z.p.p // case 1

9. **else if** z == z.p.right

10. z = z.p // case 2

11. Left-rotation (T, z) // case 2

12. z.p.color = black // case 3

13. z.p.p.color = red // case 3

14. Right-rotation (T, z.p.p) // case 3

15. **else** (same as **then** clause with "right" and "left" exchanged)

16. T.root.color = black ← may just insert the root or the red violation reach root



INSERT: case 1




z 's "uncle" (y) is **red**

Idea: (z is a right child)

- ▶ $p[p[z]]$ (z 's grandparent) must be black: $p[z]$ is red

- ▶ Color $p[z] \leftarrow$ **black**
- ▶ Color $y \leftarrow$ **black**
- ▶ Color $p[p[z]] \leftarrow$ **red**
- ▶ $z = p[p[z]]$

- Push the "**red**" violation up the tree





INSERT: case 1



z 's "uncle" (y) is **red**

Idea: (z is a left child)

- ▶ $p[p[z]]$ (z 's grandparent) must be black: $p[z]$ is red

- ▶ Color $p[z] \leftarrow$ **black**
- ▶ Color $y \leftarrow$ **black**
- ▶ Color $p[p[z]] \leftarrow$ **red**
- ▶ $z = p[p[z]]$

- Push the "**red**" violation up the tree








INSERT: case 3

Case 3:

- ▶ z's "uncle" (y) is **black**
- ▶ z is a left child

Idea:

- ▶ Color $p[z] \leftarrow$ **black**
- ▶ Color $p[p[z]] \leftarrow$ **red**
- ▶ $\text{RIGHT-ROTATE}(T, p[p[z]])$
 - ▶ No longer have 2 reds in a row
 - ▶ $p[z]$ is now black





INSERT: case 2



Case 2:

- ▶ z's "uncle" (y) is **black**
- ▶ z is a right child

Idea:

- ▶ $z \leftarrow p[z]$
- ▶ `LEFT-ROTATE(T, z)`


\Rightarrow now z is a left child, and both z and p[z] are red \Rightarrow case 3






Example


Insert 4




Case 1



Case 2



Case 3





Complexity analysis

- ▶ Time complexity of detailed steps
 - A red-black tree has $O(\log n)$ height
 - Search for insertion location takes $O(\log n)$ time
 - Addition to the node takes $O(1)$ time
 - The while loop will be executed at most $O(\log n)$ time
 - Each recoloring and each rotation take $O(1)$ time
 - Never performs more than two rotations, since the loop terminates if case 2 or case 3 is executed
 - Hence, an insertion in a red-black tree takes $O(\log n)$ time

What are the advantages of red-black tree over AVL tree?



Exercises

- ▶ What is the ratio between the longest path and the shortest path in a red-black tree?
 - The shortest path is at least $bh(\text{root})$
 - The longest path is equal to $h(\text{root})$
 - Since $h(\text{root}) \leq 2bh(\text{root})$, the ratio is ≤ 2
- ▶ When we insert a node into a red-black tree, we initially set the color of the new node to red. Why didn't we choose to set the color to black?



Red-black trees: summary

- ▶ Red-black trees guarantee that the height of the tree will be $O(\log n)$

- ▶ Operations on red-black-trees:
 - SEARCH $O(h)$
 - PREDECESSOR $O(h)$
 - SUCCESSOR $O(h)$
 - MINIMUM $O(h)$
 - MAXIMUM $O(h)$
 - INSERT $O(h)$
 - DELETE $O(h)$




Recommended reading

- ▶ Reading
 - Chapter 13, textbook
- ▶ Next lectures
 - Heap, chapters 6&12, textbook



DELETE operation




DELETE: the color of the node to be removed -- **red**

1. Every node is either **red** or **black** OK!
2. The root is **black** OK!
3. Every leaf (NIL) is **black** OK!
4. If a node is **red**, then both its children are **black** OK!
5. For each node, all paths from the node to descendant leaves contain the same number of **black** nodes OK!

Note: the deletion of a red node is the same as the deletion of a node in BST



DELETE operation



DELETE: the color of the node to be removed -- **Black**

1. Every node is either **red** or **black** OK!
2. The root is **black** Not OK! If removing the root and the child that replaces it is **red**
3. Every leaf (NIL) is **black** OK!
4. If a node is **red**, then both its children are **black** Not OK! Could create two red nodes in a row
5. For each node, all paths from the node to descendant leaves contain the same number of **black** nodes Not OK! Could change the black heights of some nodes




Deletion on red-black tree

- ▶ Similar to the deletion on BST, but need to use an auxiliary procedure **RB-Delete-Fixup** to restore the red-black tree properties
- ▶ Four different cases of **RB-Delete-Fixup**
 - Case 1: x's sibling w is red
 - Case 2: x's sibling w is black, and both of w's children are black
 - Case 3: x's sibling w is black, w's left child is red, and w's right child is black
 - Case 4: x's sibling w is black, and w's right child is red (left child either color)



Case 1


- ▶ Case 1: x 's sibling w is red
 - Solution: rotate and recolor





Case 2


- ▶ Case 2: x 's sibling w is black, and both of w 's children are black
 - Solution: recolor





Case 3

- Case 3: x 's sibling w is black, w 's left child is red, and w 's right child is black





Case 4

- Case 4: x 's sibling w is black, and w 's right child is red (left child either color)

